

AI FOR BUSINESS LEADERS · ONLINE SERIES FOR ISCA

AI Governance Overview

Ray Han, PhD (Computer Science, NTU)

Chief AI Trainer, HGT Consultancy Pte Ltd

Would you approve the launch?



HYPOTHETICAL EXAMPLE

A Singapore retailer's customer-service chatbot

The supplier says its product follows Singapore’s Model AI Governance Framework.



Approve the launch

Supplier alignment alone does not establish readiness.



Reject the proposal

The available facts do not justify permanent rejection.



Request more information

Withhold approval and obtain evidence about this deployment.

Answer: Request more information

Questions to consider first

01	Data*	What personal data will the chatbot collect or use?
02	Customer commitments	What commitments can it make to customers?
03	Jurisdictions	Which jurisdictions will it serve?
04	Accountability	Who will be accountable for its operation?

*Direct PDPA relevance where personal data is involved. Other questions may also raise data-protection obligations.

What is an AI system?

A machine-based system that infers from inputs how to produce predictions, content, recommendations or decisions, with varying autonomy, and outputs that can affect the real world.

Three types of AI you need to know.

Predictive model

Estimates credit risk.

Generative AI

Produces a customer reply.

AI agent

Acts through software tools.

What is AI governance?

The responsibilities, policies, decision processes and controls through which an organisation oversees AI it builds, buys or uses.

It answers three questions.

Who decides?

Who approves a use, on what evidence, and when must they reconsider?


What could go wrong?

What harm is possible, and which controls reduce it?

What must we meet?

Which legal and contractual obligations apply to this use?

Governance covers the AI lifecycle



Reassess whenever the model, data, users or operating context changes.

Six components of organisational AI governance

01 Accountability and oversight

Owner, approval authority and escalation.

02 Inventory and impact assessment

Uses, purposes, affected people and consequences.

03 Data and supplier governance

Privacy, quality, security and supplier responsibilities.

04 Testing and human oversight

Acceptable performance and effective intervention.

05 Transparency and recourse

Accurate explanations and usable review routes.

06 Monitoring and improvement

Outcomes, incidents, changes and retirement.


Who reviews an AI proposal before launch?

Role	Question to resolve	Evidence before approval
Business owner	What is the intended use and acceptable outcome?	Purpose, boundaries and affected people
Technical and data teams	Does the system work acceptably in this setting?	Test results, limitations and operating controls
Risk, legal and compliance	Which obligations apply, and are the safeguards adequate?	Applicability assessment and independent challenge
Accountable executive	Do the evidence and conditions justify deployment?	Approval record, residual risk and review triggers

AI governance across jurisdictions

Assess the use, the affected people, the locations served, and whether the organisation is a provider or a deployer.

Select a jurisdiction to see its approach.



BINDING AI LAW	GUIDANCE AND EXISTING LAW
European Union	Singapore
South Korea	United Kingdom
China	Japan
FRAGMENTED STATE LAW	India
United States	Malaysia

Different instruments have different legal effects

Binding requirements	Acts, regulations and binding regulatory requirements create duties for the relevant activity.
-----------------------------	--

Supervisory expectations	Guidance may explain existing law or inform supervisory assessment and corrective action.
---------------------------------	---

Voluntary frameworks	Methods and guidance for organising controls. Contracts or policies may add commitments.
-----------------------------	--

AI Governance in Singapore

Binding requirements	The PDPA governs relevant collection, use and disclosure of personal data, subject to its scope and exceptions. Other business and customer obligations may also apply.
Supervisory expectations	Sector instruments add detail where the organisation conducts a regulated activity. A financial institution reviews the relevant MAS requirements and guidance as well.
Voluntary frameworks	The Model AI Governance Framework covers internal governance, human involvement, operations management and stakeholder communication. It does not replace applicable law.

EU AI Act

Non-EU organisations can be in scope, including certain uses of AI outputs in the Union.

WHAT THE ACT REGULATES

Prohibited practices Uses the Act bans within its scope.	High-risk systems Risk management, data, records and human oversight.	General-purpose AI Separate obligations for model providers.	Transparency Disclosure for specified AI interactions and content.
--	---	--	--

WHEN EACH PART APPLIES

2 February 2025 Prohibitions	→ 2 August 2025 General-purpose AI provisions	→ 2 August 2026 General application, including Article 50	→ 2 December 2027 Annex III high-risk requirements	→ 2 August 2028 Annex I product-related high-risk requirements
--	---	---	--	--

Applicable obligations and potential harm are different questions

Which duties apply?

The obligations and supervisory expectations relevant to this use.

What harm is possible?

Who could be harmed, how severely, how likely, and whether it can be corrected.

Compare two uses in the same retailer.

Drafting a meeting summary. An employee checks it before circulating, so the effect is contained.

Automatically refusing refunds. This directly affects customers, so the decision and review arrangements need closer examination.

A voluntary framework label does not tell an approver which use presents the greater risk.

Does following a framework resolve the data issue?

The retailer collects contact details and purchase histories to personalise offers, and has not explained those purposes. The supplier follows the Model AI Governance Framework and displays an AI label. Assume the PDPA notification obligation applies.

 **Yes. Following the framework resolves it.**
Framework alignment does not establish that the retailer met its obligation.

 **No. Address the applicable PDPA obligation before launch.**
The retailer must explain the relevant data-use purposes.

 **Yes. The AI label is enough.**
An AI label alone does not explain the data-use purposes.

Answer: No. Address the applicable PDPA obligation before launch.

Foodinho: no usable route to contest

The Italian data-protection authority's notice described an automated message sent when an account was blocked. It did not tell riders how to **contest the decision** or ask for the account to be restored.

The case

Foodinho, in the Glovo group, managed delivery riders in Italy through a digital platform. Riders received an automated message when an account was blocked or deactivated.

November 2024 · the penalty

The authority imposed a five-million-euro penalty for unlawful processing of riders' data. Several infringements were involved, not one missing feature.

HSBC's credit-assessment disclosures

"From the perspective of **accountability and disclosure** of policies and practices, HSBC had acquitted itself."

Deputy Commissioner for Personal Data Protection, Singapore · paragraph 19

The case

In 2018 an unsuccessful credit-card applicant asked HSBC Singapore for its internal evaluation report. The bank supplied it with specified information redacted, and the dispute reached the Commission.

10 March 2021 · the decision

The evaluative-purpose exception supported withholding the specified data. HSBC also supplied its policies on data and AI use in credit assessment.

Consequence is broader than a regulatory fine

Neither outcome involved a regulator: one was a **customer claim**, the other a **business loss**.

Air Canada · a customer claim

A customer relied on the airline's website chatbot, which said a bereavement fare could be claimed after travel. The company's webpage said otherwise.

The tribunal rejected the argument that the chatbot was a separate entity, found negligent misrepresentation, and ordered a refund.

Zillow · an inventory loss

Zillow's 2021 filing reported a 304.4 million US dollar write-down on its home-buying business, which it then wound down.

It cited pricing unpredictability and capacity constraints, not a single defective algorithm, so the loss traces to several operating decisions.

Financial Sector in Singapore

Connect AI controls to existing governance, risk and compliance arrangements, and distinguish **binding requirements**, advisory principles and proposals rather than treating every MAS publication as having the same legal effect.

FEAT principles Advisory · first published 2018	Fairness, Ethics, Accountability and Transparency in the use of AI and data analytics.
---	--

Guidelines on AI Risk Management Consultation proposal · November 2025	Proposes an AI inventory and proportionate risk management across the lifecycle. A proposal, not an issued rule.
--	--

Financial Sector in the United Kingdom

Responsibilities follow the firm's activities and its **documented management arrangements**. Avoid assigning all AI risk to the chief technology officer, or creating a new title without decision authority.

Existing frameworks apply

The FCA identifies the Consumer Duty and senior-management accountability as relevant foundations for firms' use of AI.

Existing powers apply

The regulator can use its supervisory and enforcement powers where a firm's use of AI breaches relevant financial regulatory requirements.

Apply these governance practices in your sector

These are **questions to investigate** in your own sector, not a claim that identical duties apply. Financial-services regulation does not automatically extend to another sector.

- | | | |
|----|------------------------------|---|
| 01 | Map obligations | Start with the activity your organisation performs, and identify the requirements already attached to it. |
| 02 | Assign responsibility | An accountable owner, and an appropriate independent review. |
| 03 | Retain evidence | Testing, decisions and monitoring, so management can see whether controls are working. |

Who must own this before launch?

A UK-regulated bank is preparing a vendor-built credit-decisioning tool. Its documented management responsibilities place this lending activity with an existing senior manager. The vendor says it tests the model; the bank's risk function has not reviewed the deployment.

⊗ **The vendor. It tests and monitors the model.**

Supplier responsibilities matter and do not replace the bank's own obligations.

✓ **The responsible senior manager, with the bank's review completed first.**

Responsibility follows the firm's documented management arrangements for this activity.

⊗ **Nobody yet. An AI-specific role must be created first.**

The FCA relies on existing accountability rather than a new AI title.

Answer: The responsible senior manager, with the bank's review completed first.

Three questions, and the evidence to request

Foodinho illustrates a usable challenge route; HSBC illustrates disclosure alongside legal access limits. **Different lessons**, not one universal explainability requirement.

Question	Evidence to request
Who is accountable for this use?	A named owner, a decision authority and an escalation route
How can an affected person challenge an outcome?	The review procedure, and evidence that someone can change the decision
What have we told people about the system?	Accurate explanations of purpose, limitations and use of data

What evidence is needed before approval?

Singapore’s agentic guidance recommends human review for high-stakes or irreversible actions. This is **voluntary guidance**, not a general statutory prohibition.

01	Testing in the intended setting	Relevant users and foreseeable failure conditions, not a supplier benchmark alone.
02	A person who can intervene	With the authority, information and time to do so.
03	Production monitoring	What it detects, who receives alerts, and when escalation is required.
04	A tested way to stop	Stop, restrict or roll back — tried, not assumed.

Would you approve the chatbot now?



HYPOTHETICAL EXAMPLE

A Singapore retailer's customer-service chatbot

The team has supplied evidence. Refund-policy answers have not passed the retailer’s acceptance tests.

✓	Data	Disclosures explain what is collected and why.
✓	Customer commitments	Refund answers withheld; product and order status only.
✓	Jurisdictions	Singapore customers only in this pilot.
✓	Accountability	Head of Customer Operations owns it; a human reviewer can correct an outcome.

Restrict the pilot. Permit product information and order status. Keep refund decisions with staff until testing and controls justify expansion.

Approve, restrict or pause?

An insurance tool flags claims for manual review. It has an owner, a review procedure and accurate customer information. Testing has used only a supplier benchmark. There is no production monitoring and no rollback process.

⊗ **Approve production use.**

The stated gaps would remain unresolved.

⊗ **Run an unrestricted pilot.**

Calling a deployment a pilot does not establish the missing safeguards.

✓ **Pause production and establish the safeguards.**

A bounded pilot may follow, once testing, monitoring, human review and rollback are in place.

Answer: Pause production and establish the safeguards.

Record the next decision for each priority AI use

✓	Use case	Complaints triage · hypothetical
✓	Accountable owner	Head of Customer Operations
✓	Review route	Some escalated complaints close without human review
✓	Evidence	Supplier test data only
✓	Next decision, and by when	Whether a controlled pilot can proceed once those gaps are addressed

Which action should come first?

A firm has completed its AI inventory. Next week it plans to activate a tool that automatically declines insurance claims. There is no appeals route and no production monitoring. The board has also asked about ISO certification.

⊗ **Begin ISO certification.**

Certification may provide assurance, and it does not resolve next week's control gap.

✓ **Pause the launch and establish the missing controls.**

The known exposure is the automatic decline with no appeals route or monitoring.

⊗ **Extend the inventory to every remaining use.**

Worth doing, and it does not address this known exposure.

Answer: Pause the launch and establish the missing controls.

AI Governance implementation roadmap

LIFECYCLE · 5 STEPS

- 1. Define
- 2. Develop or procure
- 3. Test and approve
- 4. Operate and review
- 5. Retire or replace


COMPONENTS · 6

- 1. Accountability and oversight
 - 2. Inventory and impact assessment
 - 3. Data and supplier governance
- 4. Testing and human oversight
 - 5. Transparency and recourse
 - 6. Monitoring and improvement

ASK OF ANY PROPOSAL · 3

- 1. Who is accountable for this use?
- 2. How can an affected person challenge an outcome?
- 3. What have we told people about the system?

ACTION PLAN



References

1	Singapore Model AI Governance Framework · second edition, 2020
2	PDPC · Data Protection Obligations
3	OECD · Organisation for Economic Co-operation and Development · 2024 explanatory memorandum
4	NIST AI RMF 1.0
5	ISO/IEC 42001:2023
6	European Commission · AI Act overview
7	UK AI regulation white paper · 2023
8	Malaysia · National Guidelines on AI Governance and Ethics, 2024 · AI Governance Bill consultation, July 2026
9	Japan · Act on Promotion of R&D and Utilization of AI-Related Technologies, 2025
10	MAS · Supervisory approach and regulatory instruments

11	Official implementation timeline · July 2026 amendment
12	Hypothetical example · course teaching synthesis
13	Garante · Foodinho (Glovo) · November 2024 penalty notice
14	PDPC · Re HSBC Bank (Singapore) [2021] SGPDPC 3 · paragraphs 17–19
15	Moffatt v. Air Canada · 2024 BCCRT 149
16	Zillow Group · Q3 2021 results filing · 2 November 2021
17	MAS · FEAT principles · 2018, updated 7 February 2019
18	MAS P017-2025 §3.4 · consultation proposal, finance-specific
19	FCA · AI approach · February 2026
20	UK Treasury Committee · January 2026 report
21	MDDI · human accountability for autonomous AI agents · 6 May 2026